

# Workshop Part-1

By  
Prince Khanna

# Agenda

- + What and Why about Kubernetes?
- + Architecture of Kubernetes
- + Basic Concepts of Kubernetes
- + Hands On
- + Exercise



kubernetes



kubernetes

# What is Kubernetes?

**Kubernetes**, also known as K8s, is an open-source system for automating deployment, scaling, and management of containerized applications.

Kubernetes is **an orchestration tool for containerized applications**.



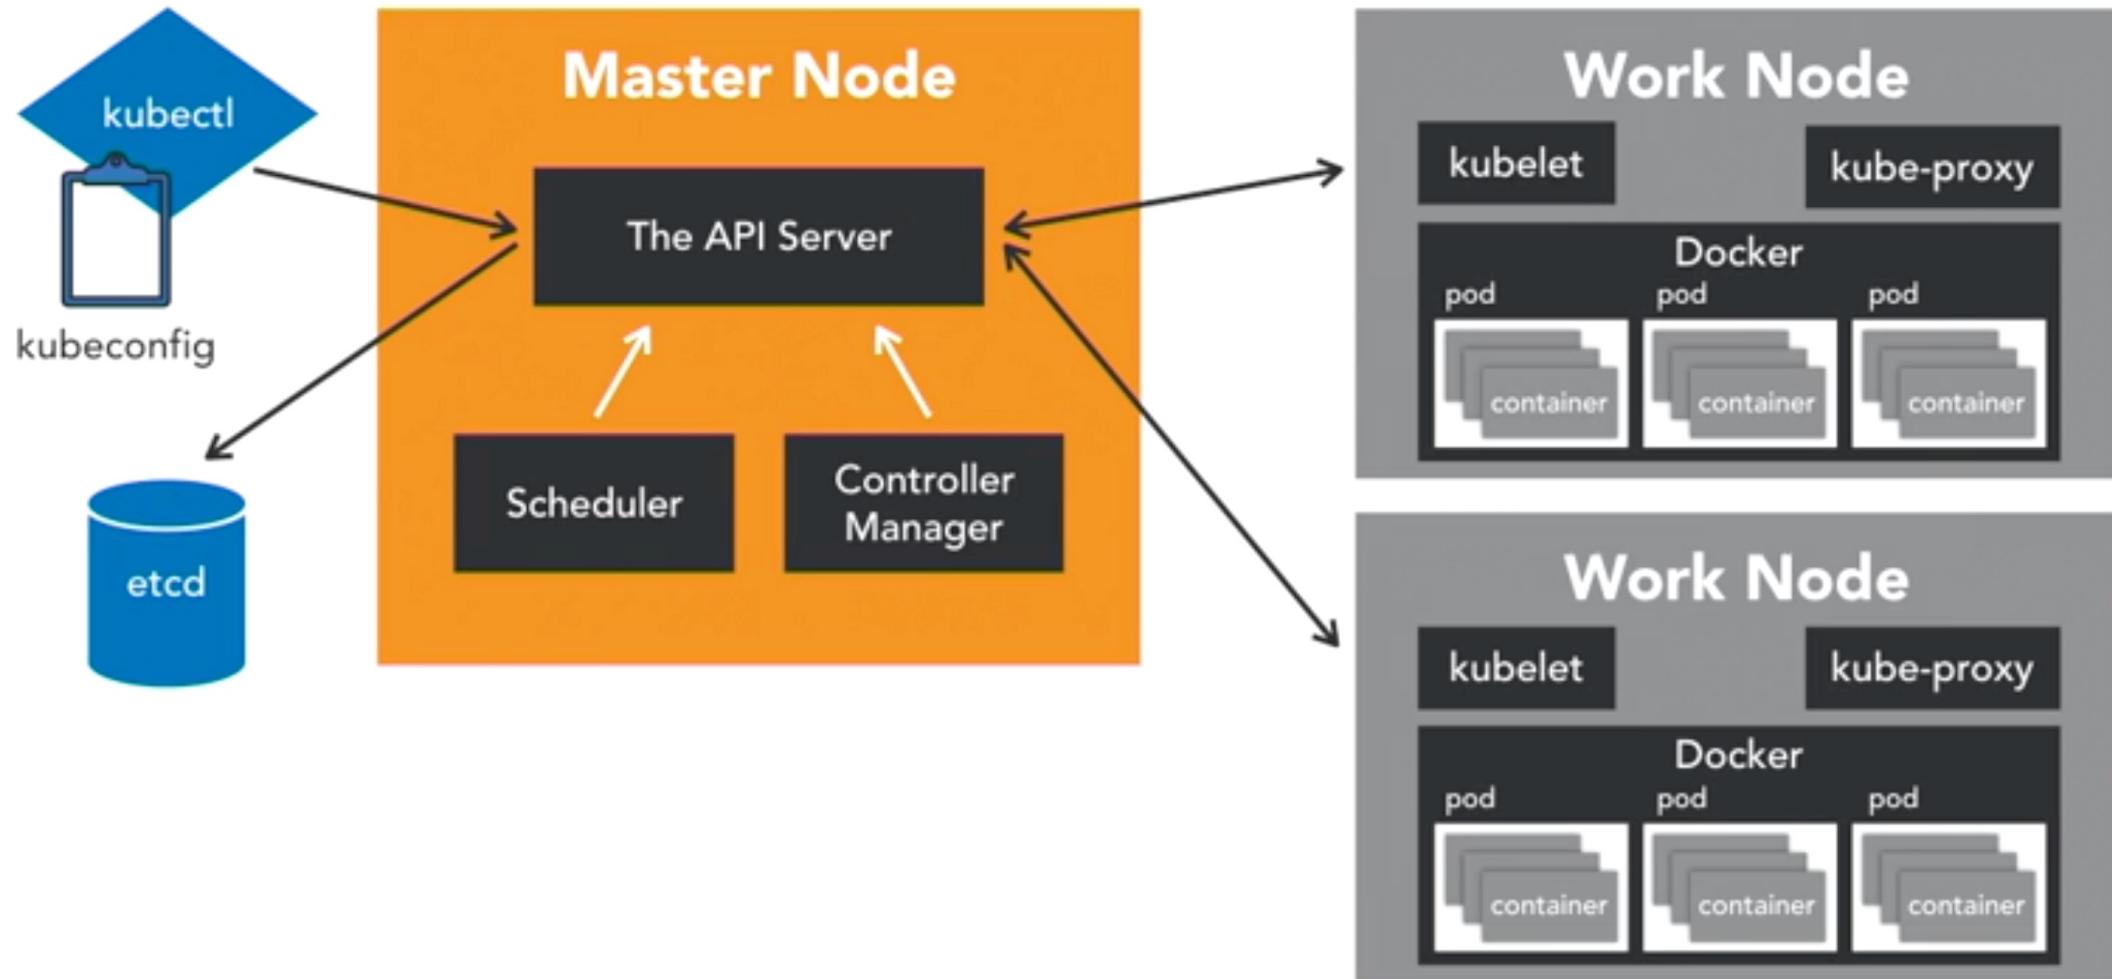
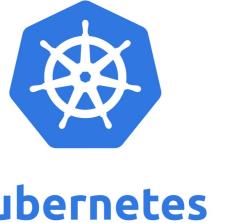
kubernetes

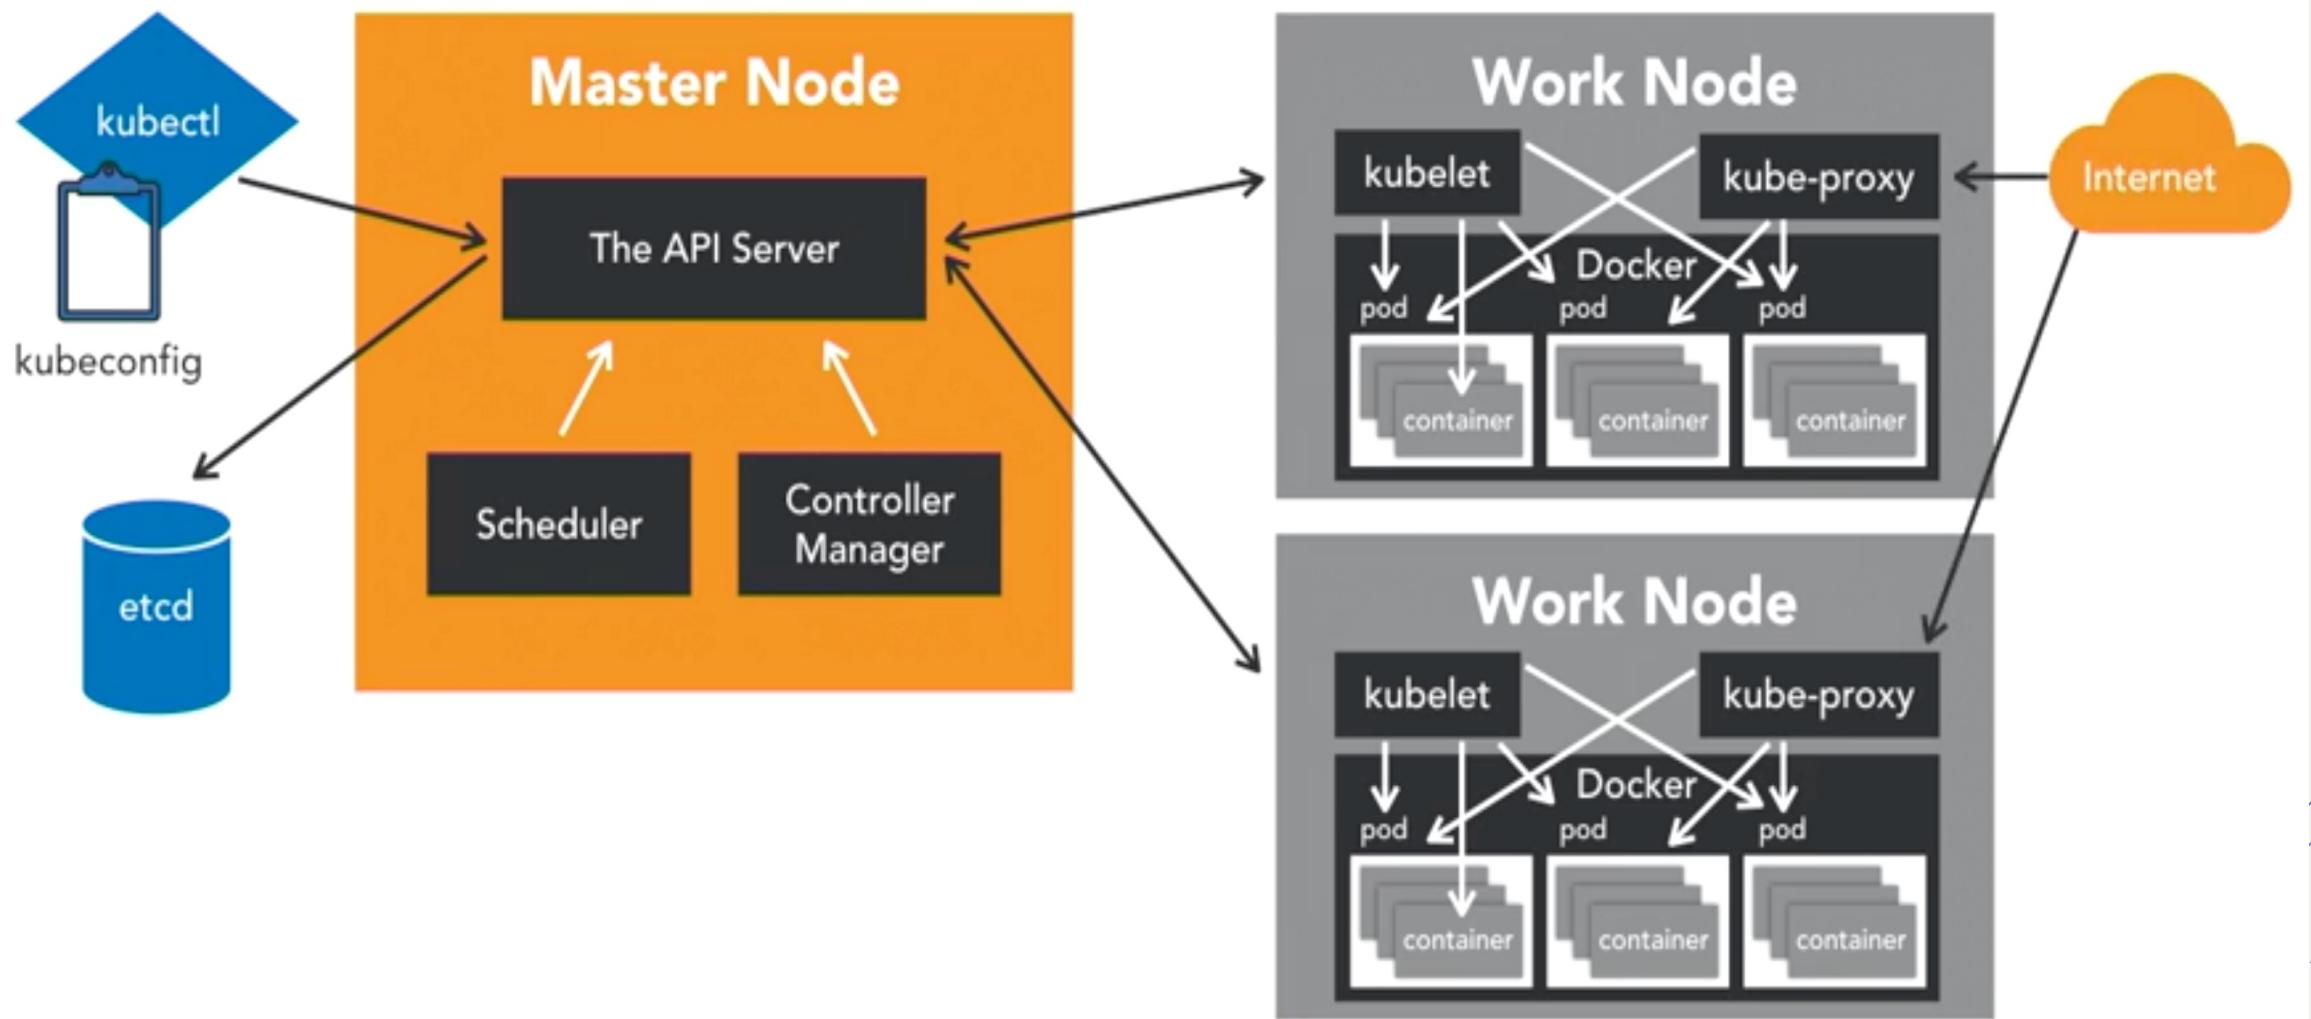
# Why Kubernetes?

Kubernetes is an orchestration tool for containerized applications. It is responsible for:

- + Deploying images and containers
- + Managing the scaling of containers and clusters
- + Resource balancing containers and clusters
- + Traffic management for services

# Architecture of Kubernetes





# Basic Concepts

- + Namespace
- + Pods
- + Replica set
- + Deployment
- + Service
- + Config Map
- + Secret
- + Volumes
- + Ingress
- + Job

# Pods

Pods are the **smallest** deployable units of **computing** that you can create and manage in Kubernetes.

A *Pod* is a group of one or more **containers**, with shared storage and network resources, and a specification for how to run the containers.

# Sample Pod yaml File

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
  ports:
  - containerPort: 80
```

# Replica set

A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical Pods.

# Deployment

A Deployment provides declarative updates for **Pods** and **ReplicaSets**.

# Sample Deployment yaml File

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

# Service

An abstract way to expose an application running on a set of **Pods** as a

Sometimes you don't need load-balancing and a single Service IP. In this case, you can create what are termed "headless" Services, by explicitly specifying "None" for the cluster IP (.spec.clusterIP). network service.

# Sample Service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

# Demo Time



# Q&A

# References

[Pods | Kubernetes](#)

[Pod Lifecycle | Kubernetes](#)

[Init Containers | Kubernetes](#)

[Deployments | Kubernetes](#)

[StatefulSets | Kubernetes](#)

[ReplicaSet | Kubernetes](#)

[TTL Controller for Finished Resources | Kubernetes](#)

[Service | Kubernetes](#)